

Glaucoma Detection Cloud Ecosystem — Project Documentation

Team: Ilia Sukhina, Yevhenii Severin, Adam Partl **Course:** Cloud Computing — Project 1

Repository: [load_balancer/](#) , [orchestrator/](#) , [cloud_ai_worker/](#) , [cloud_frontend_app/](#) , [k8s/](#) ; supporting docs [README.md](#) , [EDGE_CASES.md](#) , [LOAD_BALANCER_COMPARISON.md](#) .

Contents

1. [Project Goal and Focus](#)
 2. [Domain / Problem-Area Characterization](#)
 3. [Analysis of Similar Solutions and How Ours Differs](#)
 4. [Concept — Architecture, Views, and Technical Analysis](#)
 5. [Detailed Technical Implementation Description](#)
 6. [Experiment — Verifying the Stated Goal](#)
 7. [Conclusion and Future Work](#)
 8. [Work Distribution](#)
 9. [Appendix — Related documents](#)
-

1. Project Goal and Focus

1.1 Motivation

Glaucoma is a leading cause of irreversible blindness. Screening from retinal *fundus photographs* works but is slow: a human grader must look at every image.

The hard part is **not** the AI model — a Convolutional Neural Network (CNN) classifies one image in well under a second. The hard part is **delivering that inference reliably and elastically at clinic scale**:

- many uploads arrive at the same time;
- images vary wildly in size (a phone snapshot vs. a 4K fundus-camera export);
- a heavy TensorFlow process must not crash a web server with an Out-Of-Memory error.

So this is a **cloud-infrastructure project, not a medical-accuracy project**. We build the plumbing that runs a glaucoma CNN as a horizontally scalable, resource-aware, self-tunable service on Kubernetes.

1.2 Project focus

A **cloud-native microservice ecosystem** that:

1. Serves a glaucoma-detection CNN as an independent, statelessly scalable inference microservice.
2. Routes traffic to the *right class of worker* by request cost — small images go to a cheap CPU pool, large images go to a GPU-profile pool — using a **custom Layer-7 (HTTP) load balancer**. This is the cloud-functionality component we implement ourselves.
3. Autoscales each worker pool independently with the Kubernetes **HorizontalPodAutoscaler (HPA)**.
4. Exposes a **control-plane orchestrator** (REST API) that re-tunes routing thresholds, HPA min/max replicas and upload safeguards *at runtime* (validated, persisted), and runs a built-in synthetic load test.
5. Provides a decoupled **Streamlit dashboard** for medical staff.

1.3 Measurable objectives

The assignment asks for at least one measurable goal. We define five, all verified by one scripted experiment (`test_scalability_e2e.sh`) — see §6.

#	Objective	Pass criterion	Measured by
M1	Resource-aware routing works	100 % of small requests (< 2.5 MB) reach the CPU pool, and 100 % of large requests (≥ 2.5 MB) reach the GPU pool	<code>node_type</code> in each <code>/predict/</code> response + load-balancer logs
M2	CPU pool autoscales under load	<code>ai-worker-cpu</code> grows from 1 to ≥ 2 replicas under sustained load, then scales back down	<code>kubectl -n glaucoma get hpa/pods</code> , sampled during the test
M3	GPU pool autoscales under load	<code>ai-worker-gpu</code> grows from 1 to ≥ 2 replicas under sustained large-image load	same as M2
M4	Load is spread across the new replicas	≥ 2 distinct worker pod IDs serve requests for each pool	<code>worker_id</code> in the <code>worker_spread</code> map from <code>/api/demo/loadtest</code>
M5	Runtime reconfiguration takes effect	After <code>PUT /api/config {"size_threshold_bytes": ...}</code> , a borderline-sized request changes pool	orchestrator <code>/api/status</code> + a follow-up <code>/api/demo/predict</code> showing <code>routed_to_pool</code>

A successful run prints the verdict block shown in §6.4.

2. Domain / Problem-Area Characterization

2.1 Application domain — medical image screening

- **Input:** retinal fundus photographs (JPEG/PNG), very different in resolution and file size. The CNN expects a fixed `224 × 224 × 3` tensor, so every image is resized server-side. Result: the compute cost per request is roughly constant once the image is decoded, but the network + decode cost grows with file size — and large images are the ones that put memory pressure on a worker.
- **Output:** a probability in `[0, 1]` (likelihood of glaucoma), shown to the clinician as a percentage plus a clinical band — "Low Risk" / "Suspicious" / "High Risk".
- **Non-functional needs:** availability during screening campaigns; elasticity (load is bursty); isolation (a model crash must not take down the UI); observability (an audit trail of which node/pod handled a given medical request); and safety on untrusted uploads (size limits, type checks, decompression-bomb guards).

2.2 Cloud-computing domain — the parts we build

Cloud concept	Where it appears in this project
Containerization (Docker)	Every service ships as a Docker image; multi-stage Alpine builds keep the Go services ~10 MB.
Container orchestration (Kubernetes)	One <code>glaucoma</code> namespace; Deployments + Services; a PersistentVolume for model files; NodePort ingress for the UI and the orchestrator.
Microservices architecture	5 independently deployable services.
Elastic horizontal autoscaling	Two <code>HorizontalPodAutoscaler</code> objects driven by the Metrics Server.
Layer-7 / application-aware load balancing	Our own Go reverse proxy , routing on <code>Content-Length</code> — <i>the self-implemented cloud-functionality component</i> .
Control plane / Infrastructure-as-an-API	The Go orchestrator — a REST API that changes cluster state (HPA bounds, LB env) via <code>kubectl</code> and persists its config to a volume.
Stateless service design	Workers hold only a rebuildable model cache; any pod can serve any request.

2.3 Why resource-aware routing is the interesting cloud problem here

A plain Kubernetes `Service` load-balances **between identical replicas of one pool** (Layer 4, round-robin). It cannot look at an HTTP request and decide *which pool* should handle it.

But our workload genuinely has two cost classes — small vs. large images — and we want them on different node profiles for cost and stability reasons. Bridging that gap (a Layer-7 decision *in front of* the Layer-4 mechanics) is the functionality we implement ourselves. The algorithm is in §4.3.1, the full write-up in §5.3, and the build-vs-buy comparison in [LOAD_BALANCER_COMPARISON.md](#).

3. Analysis of Similar Solutions and How Ours Differs

3.1 Comparable approaches

The full trade-off analysis is in [LOAD_BALANCER_COMPARISON.md](#); the summary:

Approach	What it offers	Why it does not directly fit
Plain K8s Service + kube-proxy	Layer-4 load balancing across identical pods	Content-agnostic — cannot route by request body size; picks a <i>pod</i> , never a <i>pool</i> .
Kubernetes HPA alone	Elastic replica count per Deployment	Orthogonal — decides <i>how many</i> pods, never <i>which</i> pod or <i>pool</i> .
Ingress controller (NGINX / Traefik / Envoy / HAProxy)	Mature Layer 7: TLS, host/path routing, retries, rate limiting	No body-size matcher out of the box — needs custom Lua / EnvoyFilter plus a ~100 MB controller.
Service mesh (Istio / Linkerd)	Fine-grained Layer-7 traffic policy + observability	Even heavier; still no body-size matcher; overkill for 5 services.
Managed model endpoints (SageMaker / Vertex AI / Azure ML)	Managed autoscaling inference, A/B traffic splitting	Vendor lock-in; the assignment asks us to <i>implement</i> a component, not consume one; no "route by upload size" knob.
Generic API gateways (Kong / KrakenD)	Plugin-based Layer-7 routing	Request-size routing is, at best, a custom plugin — another component to run.


3.2 How our solution differs

- **The routing primitive is the differentiator.** The whole CPU/GPU decision is four lines of Go (`if r.ContentLength >= threshold { gpuProxy } else { cpuProxy }`) in an ~80-line, ~10 MB Alpine container — versus a 100 MB Ingress controller plus custom Lua/Envoy config spread across annotations, a ConfigMap and Helm values.
- **Two complementary layers, cleanly separated.** Our Go LB owns the *domain decision* (which pool); kube-proxy owns the *infrastructure mechanic* (which pod); HPA owns *how many pods*. None replaces the others — see §4.1.
- **A real control plane for the demo.** The **orchestrator** re-tunes the routing threshold, HPA min/max, target utilization and upload safeguards *at runtime* (validated, persisted, reconciled via `kubectl`), plus a built-in synthetic load tester that reports latency percentiles and per-pod spread — exactly what §6 exercises.
- **Portability.** The LB is a tiny standalone binary, so the identical routing logic runs under `docker compose` locally and under Kubernetes in the cloud deployment. An Ingress-based design would be Kubernetes-only.
- **Honest scope.** We do *not* claim mesh-grade features (TLS termination, retries, circuit breakers, rate limiting) — see §5.2. The point is to demonstrate the cloud mechanics, not to re-implement Envoy.


4. Concept — Architecture, Views, and Technical Analysis

4.1 Base architecture diagram

(a) **Request data path and topology.** The highlighted node is the self-implemented cloud-functionality component (§4.3.1, §5.3). Mermaid source: [docs/arch-a.mmd](#).



(b) Control & scaling plane (runs alongside the data path above). Mermaid source: [docs/arch-b.mmd](#).



Predict flow: clinician/operator → orchestrator (validate) → load-balancer → (size check) → CPU or GPU Service → worker pod → `model.predict` → JSON.

Responsibility split: the Go LB picks the **pool** (Layer 7) · kube-proxy picks the **pod** in that pool (Layer 4) · HPA decides **how many** pods per pool · the orchestrator **re-tunes** all of the above at runtime.


4.2 Architectural views

4.2.1 Logical / decomposition view — the five microservices

Service	Tech	What it does
Frontend	Python / Streamlit (cloud_frontend_app/app.py)	Clinician dashboard: upload a fundus image, pick a model (GET /models/), call POST /predict/ , show the probability, the clinical band, and the infrastructure telemetry (node_type , worker_id , execution_time). Fully decoupled — it knows one backend URL and does no routing.
Orchestrator	Go (orchestrator/)	Control plane / Infrastructure-as-an-API. Owns one validated, persisted JSON config. PUT /api/config reconciles the cluster (kubectl patch the HPAs, kubectl set env the LB). Runs the demo (POST /api/demo/predict , POST /api/demo/loadtest), reports GET /api/status , validates uploads <i>before</i> the LB. Algorithm: §4.3.4 .
Load balancer	Go (load_balancer/main.go)	The self-implemented cloud-functionality component — a Layer-7 reverse proxy. POST /predict/ with Content-Length ≥ SIZE_THRESHOLD → GPU backend, else → CPU backend; GET /models/ → CPU backend; GET /health → { "status": "ok" } . Logs every routing decision. Stateless, four env vars. Algorithm: §4.3.1 ; full write-up: §5.3 .
AI worker — CPU	Python / FastAPI + TensorFlow/Keras (cloud_ai_worker/main.py)	Inference microservice on the cost-effective node profile. Lazy-loads and caches Keras models, resizes the image to 224×224×3 , runs model.predict , returns the probability plus node_type , worker_id (pod hostname) and execution_time . Cheap /health (no TensorFlow load). Algorithm: §4.3.5 .
AI worker — GPU	same image as the CPU worker	Same code, deployed with NODE_TYPE="High-Performance GPU" and its own Deployment/Service/HPA, so high-resolution images land on accelerated nodes without competing with the high-throughput CPU pool.

4.2.2 Deployment / physical view

The Kubernetes objects and how they wire up are summarized below; the diagram shows the same picture with the data path (solid arrows) and the control / scaling plane (dashed arrows) drawn on top of each other. Mermaid source: [docs/deployment.mmd](#) .




Object	What it is	Why it is there
Namespace <code>glaucoma</code>	One Kubernetes namespace holding everything	On minikube it is a single node, but the design is multi-node-ready — in a real cluster the CPU and GPU Deployments would carry different <code>nodeSelector</code> / taints.
Deployments + Services	<code>ai-worker-cpu</code> , <code>ai-worker-gpu</code> , <code>load-balancer</code> , <code>frontend</code> (NodePort <code>30501</code>), <code>orchestrator</code> (NodePort <code>30900</code>)	The running services; NodePort exposes the UI and the orchestrator API outside the cluster.
PersistentVolume <code>models-pv</code> + PersistentVolumeClaim (PVC)	Holds the <code>.h5</code> / <code>.keras</code> model weights, mounted read-only into the worker pods	Model files are deliberately <i>not</i> in Git (too large) — supplied via <code>models.zip</code> + <code>prepare_models.sh</code> locally, or via PV / object storage / Git LFS in production.
ServiceAccount + role-based access control (RBAC)	A ServiceAccount for the orchestrator with <code>patch</code> rights on HPAs and Deployments in the namespace	Lets the orchestrator reconcile cluster state with <code>kubectl</code> without cluster-admin rights.
Metrics Server	Cluster add-on that reports pod CPU usage	Required by HPA; on minikube, <code>minikube addons enable metrics-server</code> .

4.2.3 Process / runtime view — the predict request lifecycle

- Upload.** The clinician uploads an image → `POST /predict/` (multipart: `file`, `model_name`) to the configured backend URL — the orchestrator, or the LB directly in the bare stack.
- Validate** (when the orchestrator is in front). Size cap → content-type allowlist → magic-byte sniff → header-only dimension check → `model_name` sanitization. The exact chain is in §4.3.4 step 4. Failures → `413` / `415` / `400`.
- Route.** The orchestrator forwards to the load balancer, which reads `Content-Length: ≥ threshold` → GPU proxy, else → CPU proxy. It logs `[route] POST /predict/ (<n> bytes) -> CPU|GPU`.
- Pick a pod.** The pool's Kubernetes Service (kube-proxy) forwards to one Ready worker pod (round-robin / random among the replicas).
- Infer.** The worker reads the bytes → resizes to `224×224`, normalizes to `[0,1]`, adds a batch dimension → `get_model(model_name)` (cache hit, or load from the PV on first use) → `model.predict(x)[0][0]` → optional synthetic CPU burn for the HPA demo → returns `{status, probability, node_type, worker_id, execution_time}`.
- Return.** The orchestrator wraps the response as `{orchestrator:{validated:{...}, routed_to_pool:"cpu|gpu"}, upstream:{...}}`; the UI shows the probability, the clinical band and the telemetry.

The same lifecycle as a UML sequence diagram — the `alt` block makes the L7 routing decision (small image → CPU pool, large image → GPU pool) explicit. Mermaid source: [docs/sequence.mmd](#).



4.2.4 Scaling / control view

- **Metrics Server** scrapes pod CPU. **HPA ai-worker-cpu** : 1 → 5 replicas at ≈70 % CPU. **HPA ai-worker-gpu** : 1 → 3 replicas at ≈75 % CPU. All four numbers are runtime-tunable through the orchestrator, within hard bounds ([0,50] replicas, [1,100] % utilization).
- Scale-up is not instant — until the new pods are `Ready`, traffic stays on the existing pods. Scale-down waits for the stabilization window. For minikube demos the workers set `LOAD_TEST_CPU_BURN_SECONDS=0.8` so inference creates measurable CPU pressure; set it to `0` for inference-only runs.
- The **orchestrator** is the human-facing control loop on top. `PUT /api/config` validates the merged config, persists it, and an `OnChange` hook runs `kubectl patch hpa ... + kubectl set env deployment/load-balancer SIZE_THRESHOLD=...`. Reconcile errors are logged at `WARN` but the API still returns the new config; `/api/config/apply` force-reconciles; `/api/config/reset` restores defaults. Failure behaviour is catalogued in [EDGE_CASES.md](#).

4.3 Detailed technical analysis of the chosen technology and its algorithms

The self-implemented component is the **Layer-7 resource-aware load balancer** (Go). Its decision is small, so we describe the "algorithm" together with the three cooperating layers.

4.3.1 The L7 routing algorithm (Go reverse proxy)

In plain words: the load balancer is a tiny HTTP reverse proxy. For each `POST /predict/` request it reads the `Content-Length` header, and — without ever touching the body — forwards the request unchanged to either the

GPU backend (if the upload is at or above a configurable byte threshold) or the CPU backend (otherwise). `GET /models/` always goes to the CPU backend; `GET /health` is answered locally. Every decision is logged. The whole core is these few lines:

```
const defaultThreshold = 2_500_000 // bytes (~2.5 MB), overridable via SIZE_THRESHOLD

mux.HandleFunc("/predict/", func(w http.ResponseWriter, r *http.Request) {
    if r.ContentLength >= int64(threshold) {
        log.Printf("[route] %s %s (%d bytes) -> GPU", r.Method, r.URL.Path, r.ContentLength)
        gpuProxy.ServeHTTP(w, r) // httputil.ReverseProxy to GPU_BACKEND
    } else {
        log.Printf("[route] %s %s (%d bytes) -> CPU", r.Method, r.URL.Path, r.ContentLength)
        cpuProxy.ServeHTTP(w, r) // httputil.ReverseProxy to CPU_BACKEND
    }
})

mux.HandleFunc("/models/", /* always */ cpuProxy.ServeHTTP)
mux.HandleFunc("/health", /* returns {"status":"ok"} */)
```

- **Why Content-Length?** It is in the request headers *before* the body is read, so the decision is constant-time with no buffering. File size is a good proxy for "a high-resolution fundus image that would put memory pressure on a CPU worker if many arrive at once" — the request we want on the GPU-profile pool.
- **httputil.ReverseProxy**. Go's standard-library proxy. We override the `Director` so the outgoing `Host` header is the backend's host (correct virtual hosting); the path and body pass through untouched. A backend failure yields a `502` — we do not retry (a deliberate, documented limitation; see §5.2).
- **Statelessness and config**. Four env vars only — `CPU_BACKEND`, `GPU_BACKEND`, `SIZE_THRESHOLD`, `PORT`. No persistent state, so the LB pod is freely restartable; the orchestrator changes `SIZE_THRESHOLD` via `kubectl set env`, which triggers a brief rolling restart of the (single-replica) LB.

4.3.2 The L4 layer — Kubernetes Service / kube-proxy

Each pool has a `ClusterIP` Service. kube-proxy programs iptables/IPVS rules that spread connections **across the Ready pods of that pool** (round-robin / random). It is content-agnostic — it never sees the HTTP body. That is exactly why the Layer-7 LB sits *in front of it*: the LB picks the **pool**, kube-proxy picks the **pod**.

4.3.3 The autoscaling control algorithm — HorizontalPodAutoscaler

In plain words: every few seconds the HPA controller looks at how busy a pool's pods are (average CPU) and resizes the pool so that busyness lands near a target. Concretely it applies the standard ratio formula:

```
desiredReplicas = ceil( currentReplicas × ( currentMetricValue / targetMetricValue ) )
```

clamped to `[minReplicas, maxReplicas]` (CPU `[1,5]` @ 70 %, GPU `[1,3]` @ 75 % by default). Scale-up is prompt but bounded by pod start time; scale-down is delayed by a stabilization window to avoid flapping. HPA is **orthogonal** to routing — it changes *how many* pods exist per pool, never *which* pod or pool a request hits. (The three layers in one line: **LB = which pool, kube-proxy = which pod, HPA = how many pods.**)

4.3.4 The control-plane algorithm — orchestrator config reconciliation

1. **Accept the change.** `PUT /api/config` takes a *partial* JSON body. Unknown fields → `422` (so typos surface). A field explicitly set to `null` is a no-op.
2. **Validate the result.** The patch is merged onto the current config, then validation runs **on the merged result** — hard bounds plus cross-field rules (e.g. `cpu_min_replicas ≤ cpu_max_replicas`; each `allowed_content_types` entry must start with `image/`). Any violation → `422`, server state unchanged.

3. **Persist, then reconcile.** On success the config is written atomically to disk, then an `OnChange` hook reconciles the cluster: `kubectl patch hpa ai-worker-cpu / ai-worker-gpu (new min/max/target)` and `kubectl set env deployment/load-balancer SIZE_THRESHOLD=<n>`. Reconcile errors are logged at `WARN`; the API still returns the now-persisted config. `POST /api/config/apply` re-runs the reconciliation (useful right after deploy, or if `kubectl` was unavailable earlier) and returns `207 Multi-Status` with a per-step error list if any patch failed.

4. **Upload safeguard chain** (on `POST /api/demo/predict`, run before forwarding to the LB):

1. size cap via a `limit+1` read — never buffers an oversized body → `413`;
2. declared content-type must be in the allowlist → `415`;
3. magic-byte sniff (`http.DetectContentType`) → `415`;
4. header-only `image.DecodeConfig` to enforce `max_image_pixels` — **the decompression-bomb guard**: a 1 KiB PNG that *claims* $100\,000 \times 100\,000$ is rejected without allocating the pixel buffer → `413`;
5. `model_name` sanitization — `[A-Za-z0-9._-]` only, no `..` segments → `400`.

5. **Synthetic load test** (`POST /api/demo/loadtest`):

- fires `n` `predict` requests at the LB with the configured `concurrency`. Both are capped — `n ≤ max_loadtest_requests`, `concurrency ≤ max_loadtest_concurrency` — else `422 before any traffic`;
- if no `image_base64` is supplied it generates a tiny synthetic JPEG (zero prerequisites);
- reports duration, throughput (requests per second), success/failure counts, latency percentiles (`min/avg/p50/p95/p99/max`), `node_distribution`, `worker_spread` (requests per worker pod — the key signal that HPA actually spread the load), and an HTTP `status_codes` breakdown;
- a client disconnect lets in-flight requests finish, launches no new ones, and returns a partial report.

4.3.5 The inference algorithm — CNN on the worker (FastAPI + TensorFlow/Keras)

Step	Detail
Model management	Lazy loading + an in-process cache (<code>MODEL_CACHE</code> dict). On the first request for a filename the worker loads it from <code>MODEL_DIR</code> (the PV) with <code>load_model(path, compile=False)</code> — inference-only, so no optimizer state, less RAM, faster init. Later requests are cache hits. Supports <code>.h5</code> , <code>.keras</code> , and TF SavedModel directories. <code>GET /models/</code> lists the available files for the UI dropdown.
Preprocessing	<code>PIL.Image.open(BytesIO(bytes)).convert("RGB")</code> (normalizes grayscale/RGBA) → <code>resize((224,224))</code> (the CNN's expected input — typical for ResNet/VGG-style backbones) → <code>img_to_array(img) / 255.0</code> (normalize to <code>[0,1]</code>) → <code>np.expand_dims(..., axis=0)</code> (final shape <code>(1, 224, 224, 3)</code>). Done entirely in memory; the upload is never written to disk.
Inference	<code>model.predict(x)</code> returns a <code>(1,1)</code> array; the scalar is the glaucoma probability ($\times 100$ for the UI). The CNN is a standard image classifier: stacked convolution + pooling feature extractor → dense layers → sigmoid. The project consumes a <i>pre-trained</i> model file; training is out of scope.
Telemetry	Every response carries <code>node_type</code> (from <code>NODE_TYPE</code> — "Standard CPU" vs "High-Performance GPU"), <code>worker_id</code> (<code>socket.gethostname()</code> = the pod name — the audit trail and the scaling proof), and <code>execution_time</code> (server-measured inference latency).
Health	<code>GET /health</code> returns <code>{status, node_type, worker_id, model_dir}</code> <i>without</i> touching TensorFlow, so readiness probes stay cheap while pods are scaling.
HPA demo aid	<code>LOAD_TEST_CPU_BURN_SECONDS</code> (default <code>0</code>) optionally spins one core with deterministic floating-point work after <code>predict</code> , so a CPU-metric HPA on minikube sees measurable pressure. In production, leave it at <code>0</code> .
Error handling	A corrupted image, a missing model, etc. are caught and returned as <code>{status:"error", message:...}</code> rather than crashing the worker.

5. Detailed Technical Implementation Description

5.1 Component inventory

Component	Language / framework	Key source files	Container
Frontend	Python 3.10+, Streamlit, Requests	cloud_frontend_app/app.py , cloud_frontend_app/requirements.txt	cloud_frontend_app/Dockerfile
Orchestrator	Go (stdlib <code>net/http</code> , <code>image</code> , <code>os/exec</code> → <code>kubectl</code>)	orchestrator/main.go (entry, flags, graceful shutdown), orchestrator/config.go (defaults/validation/persistence), orchestrator/validate.go (upload safeguards), orchestrator/k8s.go (<code>kubectl</code> wrappers), orchestrator/handlers.go , orchestrator/demo.go (load-test runner), orchestrator/server.go (routes/middleware)	orchestrator/Dockerfile
Load balancer	Go (stdlib <code>net/http/httputil</code>)	load_balancer/main.go (~80 lines of code), load_balancer/main_test.go	load_balancer/Dockerfile (multi-stage Alpine, ~10 MB)
AI worker (CPU & GPU)	Python, FastAPI/uvicorn, TensorFlow/Keras, Pillow, NumPy	cloud_ai_worker/main.py , cloud_ai_worker/requirements.txt , cloud_ai_worker/tests/test_main.py	cloud_ai_worker/Dockerfile
K8s manifests	YAML	k8s/namespace.yaml , k8s/models-pv.yaml , k8s/ai-worker-cpu.yaml , k8s/ai-worker-gpu.yaml , k8s/load-balancer.yaml , k8s/frontend.yaml , k8s/ai-worker-hpa.yaml , k8s/orchestrator.yaml (+ RBAC), k8s/deploy.sh	—
Local stack	Docker Compose	docker-compose.yml , prepare_models.sh	—
Tests / experiment	Bash + <code>go test</code> + <code>pytest</code>	test_all.sh , test_load_balancer.sh , test_routing_with_logs.sh , test_scalability.sh , test_scalability_e2e.sh	—

5.2 Limitations / boundaries of the solution

The full failure catalogue is in [EDGE_CASES.md](#) . The main boundaries:

- **The load balancer is single-replica and has no HPA of its own.** If it becomes the bottleneck it does not scale. A restart (a rolling update, or a `SIZE_THRESHOLD` change) makes it briefly unavailable — there is no `PodDisruptionBudget`. It does **no** retry / circuit-breaker (a backend failure → `502`) and **no** TLS termination.
- **The orchestrator is single-replica** (`Recreate` strategy, because it mounts a `ReadWriteOnce` (RWO) PVC for its config). It is a control-plane convenience, not a hardened gateway; if overloaded it slows API/demo calls. Its

`kubectl` -based reconciliation is best-effort — if `kubectl` is unavailable the API still serves config + demo, `kubectl_available:false` is reported, and `/api/config/apply` must be re-run later.

- **No application-level admission control or queueing.** The system does **not** emit `429 Too Many Requests`; under overload, latency grows and timeouts/errors appear. Kubernetes Services have no HTTP timeout — timeouts arise client-side, or via the orchestrator's `UpstreamTimeoutSeconds` (→ `502`).
- **Kubernetes is not load-aware at the application level.** A `Ready` pod can receive another request even while busy with many; HPA reacts to CPU%, not to in-flight request count. At `maxReplicas` no more pods are added (it is raisable at runtime, but still a hard ceiling at any moment); if the cluster runs out of CPU/RAM, new pods stay `Pending`.
- **"GPU" is a profile, not real hardware in this submission.** The GPU worker runs the same image with a different `NODE_TYPE` and its own Deployment/Service/HPA; on minikube there is no accelerator and the HPA is CPU-driven (hence the optional `LOAD_TEST_CPU_BURN_SECONDS`). The architecture is ready for real GPU nodes (`nodeSelector` /taints) — that is outside this environment.
- **No medical-accuracy claims.** We validate *cloud behaviour* (routing, scaling, tunability), not diagnostic quality. The CNN is consumed pre-trained; the weights are not in Git (delivered via `models.zip` / PV / object storage / Git LFS).
- **Single-namespace, single-node demo target (minikube).** Multi-tenancy, network policies, secrets management, persistent metrics/logging stacks, and TLS are out of scope.

5.3 The self-implemented cloud-functionality component — the Go L7 Load Balancer

This is the component the assignment asks us to describe in detail as *our own implementation of a chosen cloud functionality*: **application-aware (Layer-7) load balancing**, packaged as a Docker container and run as a Kubernetes Deployment. It turns "route by request body size to a different instance class" — something a Kubernetes Layer-4 Service cannot do, and managed Ingress controllers only do via custom plugins — into a first-class, four-line decision. The routing algorithm itself is in [§4.3.1](#); the build-vs-buy comparison is in [LOAD_BALANCER_COMPARISON.md](#).

Internal design (`load_balancer/main.go` , ~80 lines)

1. Configuration — four env vars, nothing else:

Variable	Default	Meaning
<code>CPU_BACKEND</code>	<code>http://ai-worker-cpu:8001</code>	URL of the CPU worker pool's Service.
<code>GPU_BACKEND</code>	<code>http://ai-worker-gpu:8001</code>	URL of the GPU worker pool's Service.
<code>SIZE_THRESHOLD</code>	<code>2_500_000</code>	Bytes; requests at or above this go to the GPU pool.
<code>PORT</code>	<code>8080</code>	Port to listen on.

2. Two reverse proxies. `newProxy(target)` wraps `httputil.NewSingleHostReverseProxy` and overrides the `Director` so the outgoing `Host` header equals the target host (correct virtual hosting); the path and body pass through untouched. One proxy per backend.

3. Three handlers on a `http.ServeMux` :

- `POST /predict/` — the `Content-Length` branch (GPU proxy if $\geq$ `SIZE_THRESHOLD` , else CPU proxy), with a `[route]` log line;
- `GET /models/` — always the CPU proxy (also logged);
- `GET /health` — returns `{"status":"ok"}` with `Content-Type: application/json`.

4. Serve. `http.ListenAndServe(":"+port, mux)` — no extra middleware, no persistent state.

Containerization

`load_balancer/Dockerfile` is a multi-stage build: a Go builder stage compiles a static binary; the final stage is `alpine` with just that binary → ~10 MB image, millisecond cold start. Built with `docker build -t glaucoma/load-balancer:latest ./load_balancer`.

Kubernetes integration

`k8s/load-balancer.yaml` defines a Deployment (`replicas: 1` , container port `8080` , env vars pointing at the worker Services) and a ClusterIP Service `load-balancer:8080` . The orchestrator changes `SIZE_THRESHOLD` at runtime with `kubectl set env deployment/load-balancer SIZE_THRESHOLD=<n>` , which triggers a rolling restart.

Local parity

The same container also runs under `docker-compose.yml` (host `:8000` → container `:8080`), so developers exercise the *identical* routing code without a cluster.

Observability and tests

stdout logs every decision — `[route] POST /predict/ (124987 bytes) -> CPU` , `[route] POST /predict/ (3500000 bytes) -> GPU` , `[route] GET /models/ -> CPU` — auditable via `docker compose logs -f load-balancer` or `kubectl -n glaucoma logs -f deployment/load-balancer` . Unit-tested in `load_balancer/main_test.go` : threshold logic, `/health` , env handling, and a routing benchmark.

5.4 Worker, orchestrator, frontend implementation notes

Worker — `cloud_ai_worker/main.py`

A FastAPI app; the inference algorithm is in §4.3.5. The CPU and GPU pools run **the same image** — they differ only in the `NODE_TYPE` value, and each has its own Deployment/Service/HPA. Model files are read from the `models-pv` PersistentVolume.

Endpoints: `GET /models/` (list available model files), `GET /health` (cheap, no TensorFlow), `POST /predict/` (run inference).

Environment variables:

Variable	Meaning
<code>PORT</code>	Port to listen on.
<code>MODEL_DIR</code>	Directory the model files are read from (the PV mount).
<code>NODE_TYPE</code>	Label returned in every response — "Standard CPU" or "High-Performance GPU" .
<code>TF_CPP_MIN_LOG_LEVEL</code>	TensorFlow log verbosity.
<code>LOAD_TEST_CPU_BURN_SECONDS</code>	Extra synthetic CPU work after each prediction, so the HPA demo on minikube sees real load. <code>0</code> in production.

Orchestrator — `orchestrator/`

The control-plane REST API; the reconciliation algorithm is in §4.3.4.

Endpoints: `GET /healthz` , `GET /api/config` , `PUT /api/config` , `POST /api/config/apply` , `POST /api/config/reset` , `GET /api/status` , `POST /api/demo/predict` , `POST /api/demo/loadtest` .

Command-line flags (each also has an env-var form):

Flag / env	Default	Meaning
<code>-listen</code> / <code>ORCH_LISTEN</code>	<code>:9000</code>	Address to listen on.
<code>-config</code> / <code>ORCH_CONFIG</code>	<code>/var/lib/orchestrator/config.json</code>	Where the config JSON is persisted.
<code>-lb-url</code> / <code>ORCH_LB_URL</code>	<code>http://load-balancer:8080</code>	Load-balancer URL the demo sends traffic to.
<code>-namespace</code> / <code>ORCH_NAMESPACE</code>	<code>glaucoma</code>	Kubernetes namespace it reconciles.
<code>-kubectl</code> / <code>ORCH_KUBECTL</code>	<code>kubectl</code>	Path to the <code>kubectl</code> binary.
<code>-no-kube</code> / <code>ORCH_NO_KUBE</code>	<code>off</code>	Config-only mode — serves the API without touching the cluster (for local runs).
<code>-read-timeout</code>	<code>30s</code>	HTTP read timeout.
<code>-write-timeout</code>	<code>120s</code>	HTTP write timeout (long, so load tests can finish).

Config fields, allowed range, and default:

Field	Range	Default
<code>size_threshold_bytes</code> — CPU↔GPU routing threshold	1 KiB ... 1 GiB	2 500 000
<code>cpu_min_replicas</code> / <code>cpu_max_replicas</code>	0 ... 50	1 / 5
<code>gpu_min_replicas</code> / <code>gpu_max_replicas</code>	0 ... 50	1 / 3
<code>cpu_target_utilization</code> / <code>gpu_target_utilization</code>	1 ... 100 %	70 / 75
<code>max_upload_bytes</code> — largest accepted upload	4 KiB ... 64 MiB	16 MiB
<code>max_image_pixels</code> — decompression-bomb guard	4 096 ... 64 000 000	16 000 000
<code>allowed_content_types</code>	each must start with <code>image/</code>	<code>image/jpeg</code> , <code>image/png</code>
<code>max_loadtest_requests</code>	1 ... 5 000	200
<code>max_loadtest_concurrency</code>	1 ... 256	16
<code>upstream_timeout_seconds</code>	1 ... 600	60

Frontend — `cloud_frontend_app/app.py`

A Streamlit dashboard. It reads one env var, `BACKEND_URL`. On load it calls `GET /models/` to fill the model dropdown; on upload it calls `POST /predict/` and shows the probability, the clinical band (Low Risk / Suspicious / High Risk) and the infrastructure telemetry (`node_type`, `worker_id`, `execution_time`). It is stateless and does no routing itself — it only knows the one backend URL.

5.5 Step-by-step: configuring and running the whole system

Option A — Local development (Docker Compose)

```
# 0. Prerequisites: Docker + Docker Compose.
# 1. Provide model weights (not in Git): put models.zip in the repo root, then
./prepare_models.sh           # extracts models.zip → ./models
# 2. Build and start the whole stack (frontend, load balancer, CPU worker, GPU worker)
docker compose up --build
#   Frontend:      http://localhost:8501
#   Load Balancer: http://localhost:8000   (GET /health, POST /predict/, GET /models/)
# 3. Watch routing decisions live
docker compose logs -f load-balancer
# 4. Verify routing (generates a ~500 KB and a ~3.5 MB image)
docker compose up --build -d && ./test_load_balancer.sh
./test_routing_with_logs.sh http://localhost:8000 compose   # several sizes + matching LB logs
```

Option B — Cloud deployment (Kubernetes / minikube)

```
# 0. Prerequisites: Docker, minikube, kubectl (or: alias kubectl="minikube kubectl --").
# 1. Start the cluster and the Metrics Server (required for HPA)
minikube start --driver=docker
minikube addons enable metrics-server
# 2. Provide model weights as in Option A.
# 3. Build all images into minikube's Docker and apply every manifest in k8s/
./k8s/deploy.sh           # build + apply; later: ./k8s/deploy.sh apply (re-apply only)
#   Creates in namespace 'glaucoma': namespace, models-pv (PV+PVC), ai-worker-cpu/-gpu (Deploy+Svc),
#   load-balancer (Deploy+Svc), frontend (Deploy+Svc NodePort 30501),
#   ai-worker-hpa (cpu 1→5@70%, gpu 1→3@75%), orchestrator (Deploy+Svc NodePort 30900 + RBAC)
# 4. Get the URLs and sanity-check
minikube service frontend      -n glaucoma --url
ORCH_URL=$(minikube service orchestrator -n glaucoma --url)
kubectl -n glaucoma get pods; kubectl -n glaucoma get hpa; kubectl top nodes
# 5. Inspect / tune the live config through the orchestrator
curl -s "$ORCH_URL/api/config" | jq
curl -X PUT "$ORCH_URL/api/config" -H 'Content-Type: application/json' \
  -d '{"size_threshold_bytes": 1500000, "cpu_max_replicas": 8, "gpu_max_replicas": 4}'
curl -s "$ORCH_URL/api/status" | jq '.cluster'
# (if kubectl was unavailable on a PUT:) curl -X POST "$ORCH_URL/api/config/apply"
# 6. A single validated prediction (rejects oversize / wrong type / decompression bomb / bad model_name)
curl -X POST "$ORCH_URL/api/demo/predict" -F file=@models/sample.jpg -F model_name=cnn_model_v1.keras
# 7. A synthetic batch load test (latency percentiles + per-pod spread)
curl -X POST "$ORCH_URL/api/demo/loadtest" -H 'Content-Type: application/json' \
  -d '{"n": 100, "concurrency": 8, "model_name": "cnn_model_v1.keras"}'
```

Option C — Orchestrator alone, no cluster (config-only)

```
cd orchestrator && go build -o orchestrator .
./orchestrator -listen :9000 -lb-url http://localhost:8000 -no-kube   # serves config + demo, no kubectl
```

Running the unit tests

```
cd cloud_ai_worker && pip install -r requirements.txt && python -m pytest tests/ -v # worker
cd load_balancer && go test -v ./... # LB
./test_all.sh # all + integration if k8s
```

Troubleshooting (full list in [README.md](#)): minikube /var full → `docker system prune -a --volumes`, then `minikube delete && minikube start --driver=docker --cpus=2 --memory=6000`; Metrics Server stuck at 0/1 → wait 1–2 min, or re-enable the addon; `pip install` DNS failure during the image build → set Docker DNS, or build with `--network=host`, then `./k8s/deploy.sh apply`.

6. Experiment — Verifying the Stated Goal

6.1 Purpose

Show, with a single scripted run, that the system fulfils M1–M5 ([§1.3](#)):

- **(a)** the Layer-7 load balancer routes light vs. heavy inference traffic to different worker pools;
- **(b)** Kubernetes HPA elastically adds replicas to each pool under sustained load, with traffic spread across the new pods;
- **(c)** the routing threshold is tunable at runtime via the orchestrator.

A cloud-infrastructure experiment, not a medical-accuracy test.

6.2 Setup

1. Deploy to minikube as in [§5.5 Option B](#). The worker manifests ship with `LOAD_TEST_CPU_BURN_SECONDS=0.8` so inference produces measurable CPU pressure for the CPU-based HPA on minikube.
2. Keep a port-forward to the LB open: `kubectl -n glaucoma port-forward svc/load-balancer 8000:8080`.
3. In separate terminals, watch the cluster react: `kubectl -n glaucoma get hpa -w` and `kubectl -n glaucoma get pods -w`.

6.3 Procedure

```
# args: <LB-url> <duration-seconds> <cpu-concurrency> <gpu-concurrency>
./test_scalability_e2e.sh http://localhost:8000 360 20 10
```

It sends both CPU-routed (small) and GPU-routed (large) traffic, monitors both HPAs, samples pod counts, and prints an automatic verdict.

Optional single-pool runs: `./test_scalability.sh http://localhost:8000 300 12 cpu` and ... `300 8 gpu`. For M5: `curl -X PUT $ORCH_URL/api/config -d '{"size_threshold_bytes":1500000}'`, then a `/api/demo/predict` checking `orchestrator.routed_to_pool`, plus `curl -s $ORCH_URL/api/status | jq '.cluster.hpas'`.

6.4 Expected result (success criterion)

```

CPU routing:          PASS      ← M1: every small request reached the CPU pool
GPU routing:          PASS      ← M1: every large request reached the GPU pool
CPU HPA scaling:      PASS      ← M2: ai-worker-cpu went 1 → ≥2 replicas under load
GPU HPA scaling:      PASS      ← M3: ai-worker-gpu went 1 → ≥2 replicas under load
CPU unique worker pods: 2        ← M4: load was spread over ≥2 CPU pods
GPU unique worker pods: 2        ← M4: load was spread over ≥2 GPU pods
PASS: both CPU and GPU worker pools demonstrated HPA scaling.

```

Observed run

`./test_scalability_e2e.sh http://localhost:8000 360 20 10` on minikube (Docker driver), model `cnn_model_v1.keras`. The script finished early on both pools as soon as routing + scale-up were proven, and printed exactly the verdict block above. Concrete figures from this run:

	CPU pool (ai-worker-cpu)	GPU pool (ai-worker-gpu)
Test image size	1 456 bytes (< 2.5 MB → routed to CPU)	5 889 693 bytes (≈5.9 MB, ≥ 2.5 MB → routed to GPU)
HPA CPU load observed	322 % ... 352 % / 70 % target	186 % ... 196 % / 75 % target
Replicas	1 → 2 (first new pod Ready ~24 s into the run), then → 4 while the GPU phase ran	1 → 2 (new pod pl7bs Ready within ~20 s of being created)
Requests served / status	40 (38 success , 2 transient parse-error), all node_type = Standard CPU	130, all success , all node_type = High-Performance GPU
Unique worker pods (M4)	2 — ...-pnthg : 24 reqs, ...-gr8dp : 14 reqs	2 — ...-lbgz2 : 126 reqs, ...-pl7bs : 4 reqs
Routing / Scaling verdict	PASS / PASS	PASS / PASS

`kubectl -n glaucoma get hpa mid-run` (during the GPU phase, while the CPU pool was still hot and the GPU pool ramping up):

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS
ai-worker-cpu	Deployment/ai-worker-cpu	cpu: 352%/70%	1	5	4
ai-worker-gpu	Deployment/ai-worker-gpu	cpu: 186%/75%	1	3	1 → 2

After traffic stopped, both pools cooled down (cpu: 0%) and scaled back toward their minimums once the HPA stabilization window elapsed.

Corroborating evidence:

- the load-balancer log: `[route] POST /predict/ (124987 bytes) -> CPU / (3500000 bytes) -> GPU / GET /models/ -> CPU`;
- `kubectl -n glaucoma get hpa` showing TARGETS above the threshold and REPLICAS rising during the run, then settling after the stabilization window once load stops;
- the orchestrator's `/api/demo/loadtest` response — `latency_ms` percentiles, `throughput_rps`, and a `worker_spread` map with one entry per worker pod (the clearest single signal that HPA-added replicas actually took traffic).

6.5 Interpretation

A **PASS** on all five verdict lines means the goal is met: heterogeneous inference traffic is **routed by content (size)** to the appropriate worker class by our self-implemented Layer-7 load balancer; each worker class **scales elastically** under load via Kubernetes HPA, with traffic spread across the new pods; and the routing/scaling parameters are **re-tunable at runtime** through the orchestrator control plane — all within a containerized, Kubernetes-orchestrated microservice ecosystem.

7. Conclusion and Future Work

The project delivers a small but complete cloud-native ecosystem around a glaucoma-detection CNN: five containerized microservices on Kubernetes, a self-implemented Layer-7 load balancer that routes by request size, independent horizontal autoscaling per worker pool, and a control-plane orchestrator that re-tunes the system at runtime and ships its own load test. The scripted experiment ([§6](#)) verifies all five measurable objectives.

Natural next steps, were this taken beyond a course submission:

- **Real GPU nodes** — schedule the GPU pool onto accelerated nodes via `nodeSelector` / taints instead of treating "GPU" as a label.
 - **Harden the data path** — add TLS termination, retries / circuit-breaking, and a `PodDisruptionBudget`; run the load balancer with multiple replicas.
 - **Smarter autoscaling** — drive the HPA on a custom metric (in-flight request count or queue depth) rather than CPU%, which matches the actual bottleneck more closely.
 - **Admission control** — return `429 Too Many Requests` and/or queue under overload instead of letting latency grow unbounded.
 - **Observability stack** — persistent metrics/logging (Prometheus + Grafana, a log aggregator) instead of `kubectl logs`.
 - **Multi-tenancy and secrets** — network policies, per-tenant namespaces, and a secrets manager for a real clinical deployment.
-

8. Work Distribution

The project was split along the three architectural layers introduced in [§4](#): the **inference layer** (workers, frontend, model integration), the **infrastructure layer** (the Go load balancer, all Kubernetes manifests, the orchestrator control plane), and the **verification layer** (the test framework, the HPA / scalability experiments, the edge-case catalogue). One member owned each layer end-to-end, with cross-review on the boundaries.

Member	Owned scope	Concrete deliverables
Ilia Sukhina (@sahurai)	Inference layer + project documentation. AI worker code, clinician-facing frontend, the local Docker Compose stack, all top-level written deliverables.	cloud_ai_worker/ — FastAPI service, lazy model cache, preprocessing & inference pipeline, telemetry (§4.3.5) · cloud_frontend_app/ — Streamlit dashboard, model dropdown, clinical-band rendering · baseline docker-compose.yml · README.md , LOAD_BALANCER_COMPARISON.md , DOCUMENTATION.md , architecture diagrams (docs/arch-a.png , docs/arch-b.png), sequence + deployment diagrams (docs/sequence.mmd , docs/deployment.mmd).
Adam Partl	Infrastructure layer. The self-implemented cloud-functionality component (the Go L7 load balancer), all Kubernetes manifests, and the Go orchestrator control plane.	load_balancer/main.go + load_balancer/Dockerfile — the ~80-line L7 reverse proxy with the Content-Length routing primitive (§4.3.1, §5.3) · all manifests under k8s/ — namespace, models-pv.yaml , ai-worker-cpu.yaml , ai-worker-gpu.yaml , load-balancer.yaml , frontend.yaml , orchestrator.yaml (+ ServiceAccount/RBAC), deploy.sh · the full orchestrator/ Go service — REST API, validated/persisted config, kubectl-based reconciliation, upload safeguards, synthetic load tester (§4.3.4) · per-component READMEs (load_balancer/README.md , k8s/README.md , orchestrator/README.md).
Yevhenii Severin	Verification layer. Test framework across all components, the Kubernetes autoscaling configuration, the E2E scalability experiment, and the edge-case catalogue.	Unit tests — cloud_ai_worker/tests/test_main.py + cloud_ai_worker/pytest.ini , load_balancer/main_test.go · integration / orchestration — test_all.sh , setup_k8s.sh , prepare_models.sh · HPA configuration k8s/ai-worker-hpa.yaml and the related manifest updates that made the workers HPA-friendly (resource requests/limits, <code>LOAD_TEST_CPU_BURN_SECONDS</code>) · the scalability experiment used in §6 — test_scalability.sh , test_scalability_e2e.sh , test_routing_with_logs.sh · the EDGE_CASES.md failure catalogue (committed by Ilia on Yevhenii's behalf — see commit <code>fb68219</code>).

Effort breakdown by area

The matrix below maps the work areas onto the team. **P** = primary owner (designed and implemented), **C** = contributor (reviews, follow-up patches, integration glue), blank = no involvement.

Area	Ilia	Adam	Yevhenii
AI worker (FastAPI + TF/Keras inference)	P		C
Streamlit frontend	P		
Docker Compose (local stack)	P	C	
Custom Go Layer-7 load balancer		P	C
Kubernetes manifests (Deploy/Svc/PV/RBAC)		P	C
HorizontalPodAutoscaler (HPA) setup		C	P
Orchestrator (Go REST control plane)		P	
Unit + integration test framework			P
Scalability experiment (§6)			P
Architecture, sequence, deployment diagrams	P		
Project documentation (this file, README, LB comparison)	P		
Edge-case catalogue (EDGE_CASES.md)	C		P

Authorship can be verified against the git history (`git log --pretty=format:"%h %an %s"`); the responsibility split above matches the commit attributions, with cross-contributions reflected as **C**.

Appendix — Related documents

- [README.md](#) — quick start, full command reference, troubleshooting.
- [LOAD_BALANCER_COMPARISON.md](#) — "custom Go LB vs. Kubernetes-only" trade-off analysis.
- [EDGE_CASES.md](#) — system behaviour under load, pod failures, routing/validation edge cases, network anomalies.
- [orchestrator/README.md](#) — full orchestrator API reference, config bounds, edge cases, flags.
- [load_balancer/README.md](#), [cloud_ai_worker/README.md](#), [cloud_frontend_app/README.md](#), [k8s/README.md](#) — per-component docs.